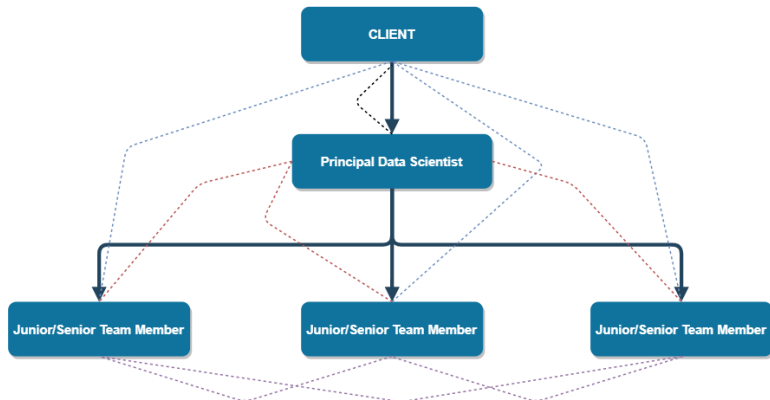



Project Management

Ryan Miller

Organizational Structure

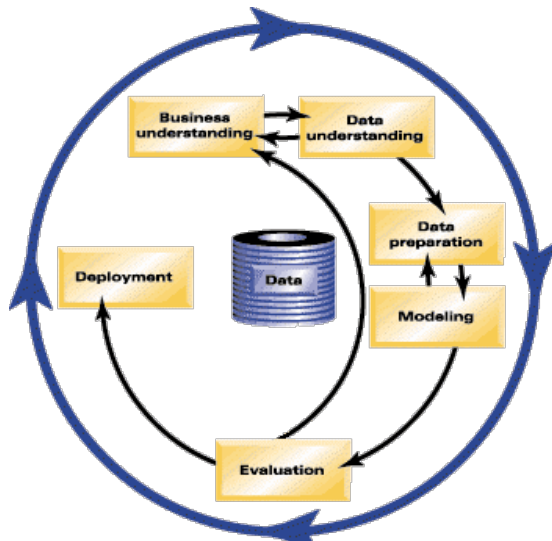


Data Science “pipelines” and “life-cycles”

- ▶ The day-to-day data science tasks take place within the organizational framework
- ▶ While “the data science process” is often discussed, its definitions vary
 - ▶ You've seen some examples at this link

CRISP-DM

We'll adopt the Cross-Industry Standard Process for Data Mining (CRISP-DM) as our model:



Phase 1 - Business Understanding

Tasks:

1. Gather necessary background information
2. Document specific specific objectives
3. Determine success criteria for the project

Each of these tasks should be undertaken in coordination with the client

Phase 1 - Objective vs. Subjective Success Criteria

- ▶ You may choose to have a mixture of objective and subjective success criteria
 - ▶ Objective = “Increase the time visitors spend on the landing page by 10%”
 - ▶ Subjective = “Identify customer clusters for targeted marketing”

Phase 2 - Data Understanding

Tasks:

1. Describing the data
2. Exploring the data
3. Verifying data quality

These tasks should be carried out at the team level (and cross-referenced with the principal and client if necessary)

Phase 3 - Data Preparation

Tasks:

- ▶ Merging/joining (ie: `left_join`)
- ▶ Selecting relevant subsets (ie: `filter`)
- ▶ Aggregating records (ie: `group_by` and `summarize`)
- ▶ Deriving new attributes (ie: `mutate`)
- ▶ Handling missing data (ie: `complete.cases` or `knnImput/rfImpute`)

These tasks should be carried out at the team level and cross-referenced by the principal (they are seldom relevant to the client at this point)

Phase 4 - Modeling

Tasks:

- ▶ Selecting a model
- ▶ Evaluating the “goodness” of a model
- ▶ Building the model
- ▶ Note: you may replace “model” with “product” in some applications

This phase is highly non-linear, it should be carried out at the team level and cross-referenced by the principal (and possibly with the client depending on their level of technical proficiency)

Phase 5 - Evaluation

Tasks:

1. Consider your model/product in regards to the business success criteria you came up with in Phase 1
2. Formalize your findings

These tasks should be undertaken in coordination with the client and principal

Phase 6 - Deployment

Tasks:

1. Deliver your model/product to the client
2. Complete wrap-up tasks (ie: technical report, etc.)

Review #1

For the following scenario determine:

1. Which phase the described actions fall under
2. Where you'd go next (and why)

A project is using medical records to build a model to predict A1c levels using more readily available measures such as blood pressure, age, weight, and waist circumference. Using the `is.na` function in R, it is discovered that 88% of the available records do not have an A1c measurement.

Review #1 - Possible Answers



- ▶ This scenario occurred during the **Data Understanding** phase
- ▶ It is possible that the 12% of medical records with A1c provides a reasonable dataset, in this case the next step would be Data Preparation (filtering out the missing data and preparing the other variables)
- ▶ It is also possible that relying on only 12% of the available records is infeasible or will induce bias into the analysis, in this case the next step would be Business Understanding (reviewing the original goals and making adjustments)

Review #2

For the following scenario determine:

1. Which phase the described actions fall under
2. Where you'd go next (and why)

In the aforementioned project you fit a linear regression model containing several variables, you receive the following model diagnostics from your software



Review #2 - Possible Answers

- ▶ This scenario occurred during the **Modeling phase**
- ▶ Linear regression doesn't appear to be an appropriate model based upon these diagnostics. It seems that a quadratic effect is being missed (hence the large negative residuals for high/low fitted values). The QQ-plot also calls into question whether the residuals are normally distributed (not a disaster for model fitting, but a problem for statistical inference)
 - ▶ For these reasons the logical next step is to return to the Data Preparation phase and explore variable transformations that might address these issues

Review #3

For the following scenario determine:

1. Which phase the described actions fall under
2. Where you'd go next (and why)

After revising the linear regression model in the previous example, a final model is chosen and is applied a "test set" of 100 new records that occurred after the original dataset was finalized. The model predicts A1c within 10% of the actual value for 86% of these new records.

Review #3 - Possible Answers

- ▶ This scenario occurred in the **Evaluation phase**
- ▶ Where to go next *depends upon the project's business goals*. If predicting A1c within 10% of the actual value for 86% of cases satisfies the previously established goals, the model is suitable for the Deployment phase. Otherwise the project might need to return to square one.